

Luca Parisi, EPCC, The University of Edinburgh
I.parisi@epcc.ed.ac.uk

11 Jun 2024

www.archer2.ac.uk





Outline



- Introduction to sampling
- Hotspot analysis
- MPI communication
- Memory usage
- Shared memory with OpenMP



Sampling vs tracing



Sampling	Tracing
Interrupts execution at constant frequency	Tracks entering/exiting function calls
Statistical representation of program execution. Might miss some very quick function calls.	All functions are correctly represented
Usually requires relinking. This can sometimes be done automatically	Often requires recompiling

- MAP is a sampling only library
- Tracing is possible using other tools like CRAY-PAT, ITAC, SCOREP ...

MAP



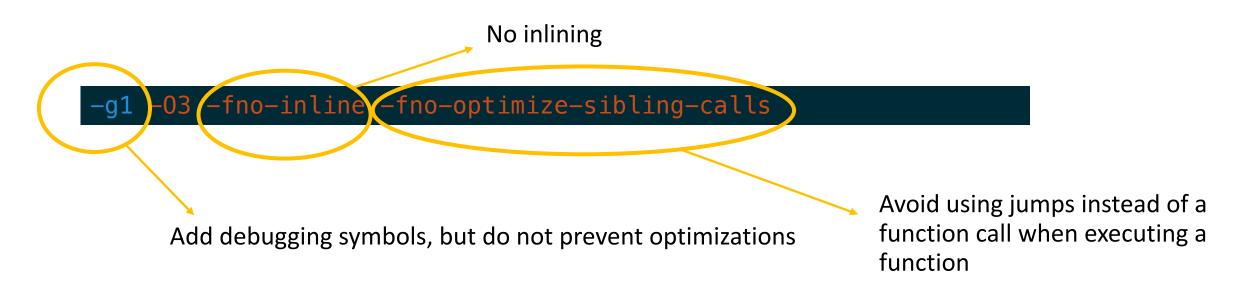
- ✓ MAP is a sampling library
- X MAP is not a tracing library. Can use other tools like CRAY-PAT, ITAC, SCOREP ...
- ✓ MAP is very easy to use and a very simple and intuitive interface
- X MAP can be imprecise
- Map lacks some more advanced features (PAPI counters, memory allocations)
- ✓ MAP gives a very quick and clear overview of performance bottlenecks



Compiling



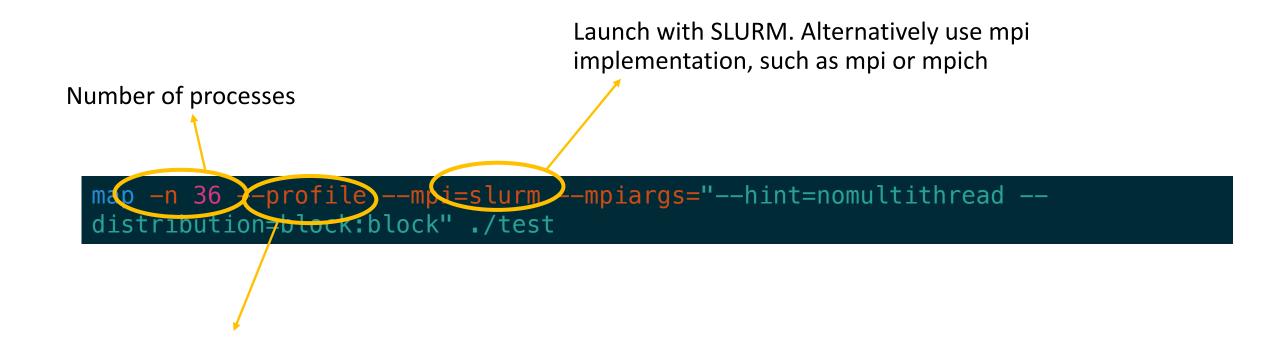
- Usually only needs to compile with debugging symbols and optimization turned on.
- Inlining can confuse the profiler
- Tail recursion optimization can confuse the stack record



Launching map



Launch from a submission script



EPCC, The University of Edinburgh

Profile from CLI, without opening the GUI

The .map files



- After execution the profile is saved as a .map file
- .map files contain links to object files and source codes. If you move, or delete, source or object file you might loose information
- You can open the .map file from the forge GUI
- You can export a .map file to a .json file

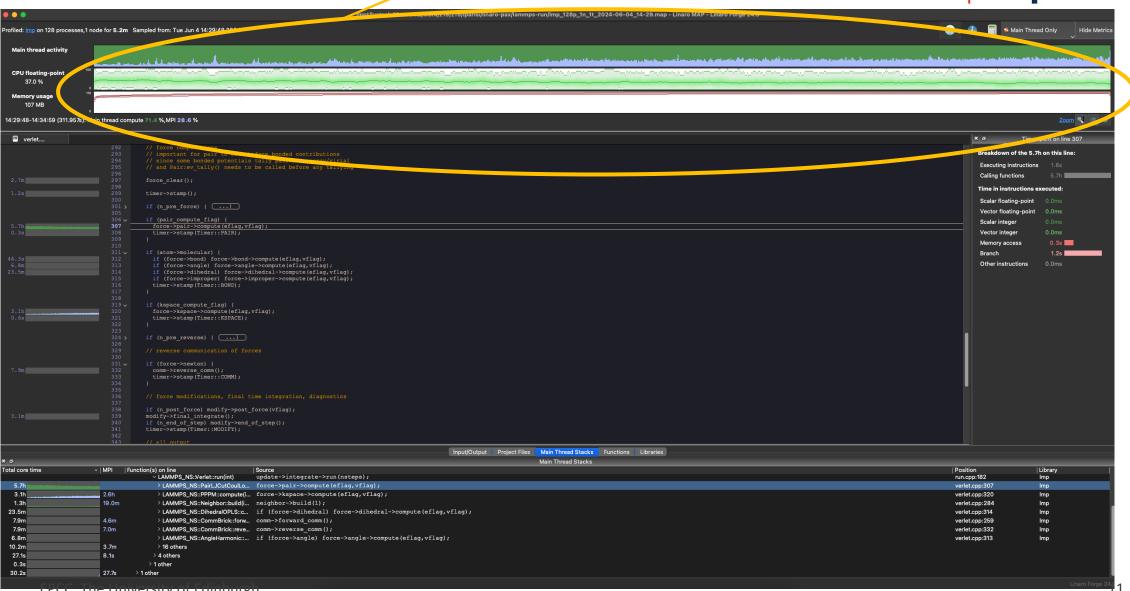
map --export=report.json resample_2p_2n_1t_2024-06-10_13-39.map



Metrics in function of time

Hotspots

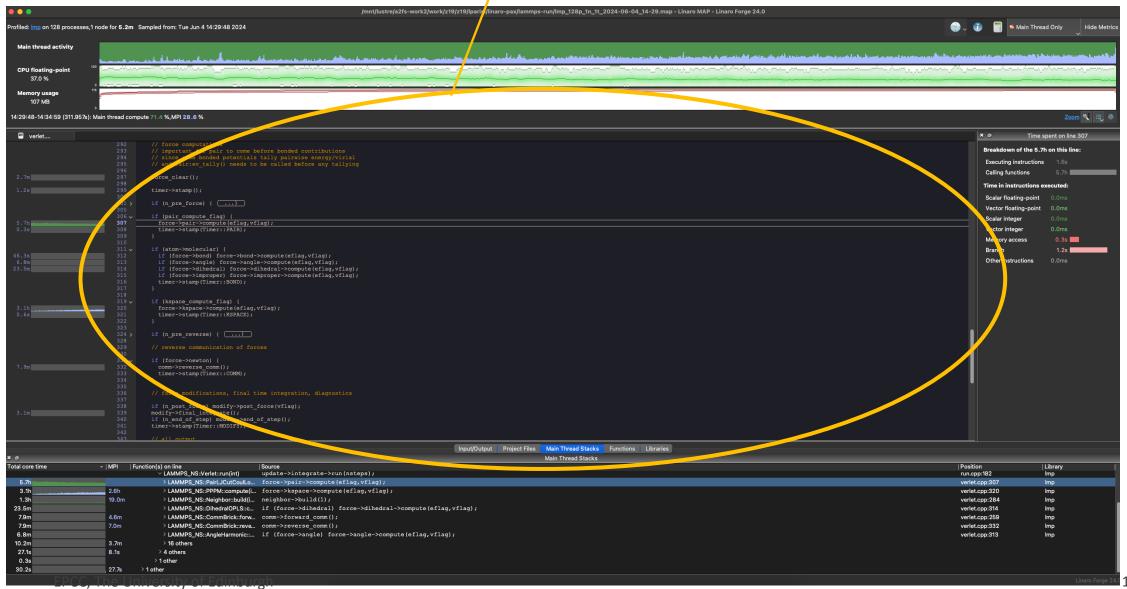




Source code

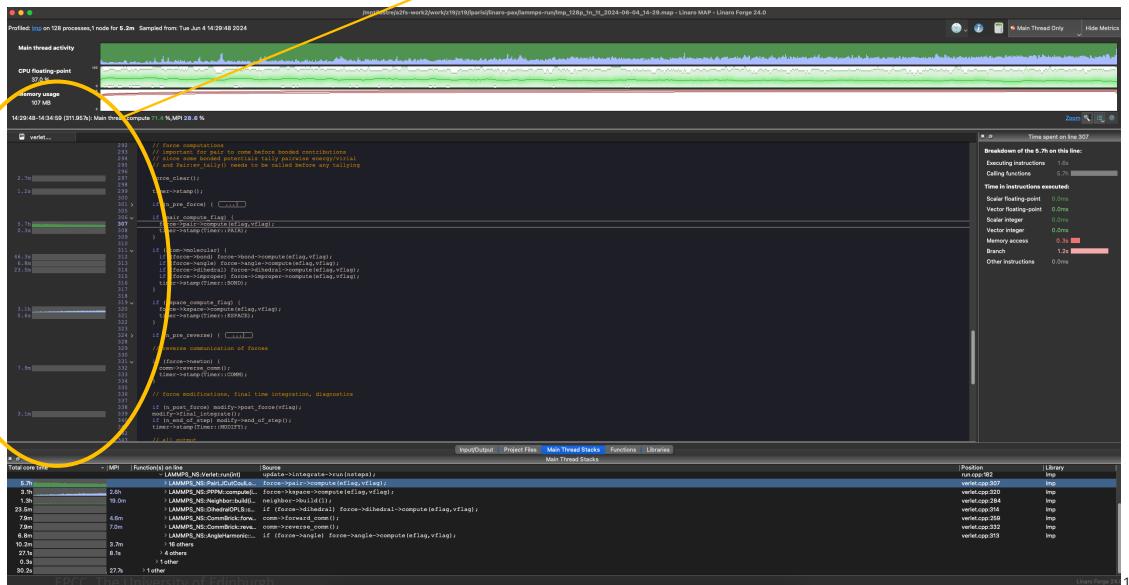
Hotspots





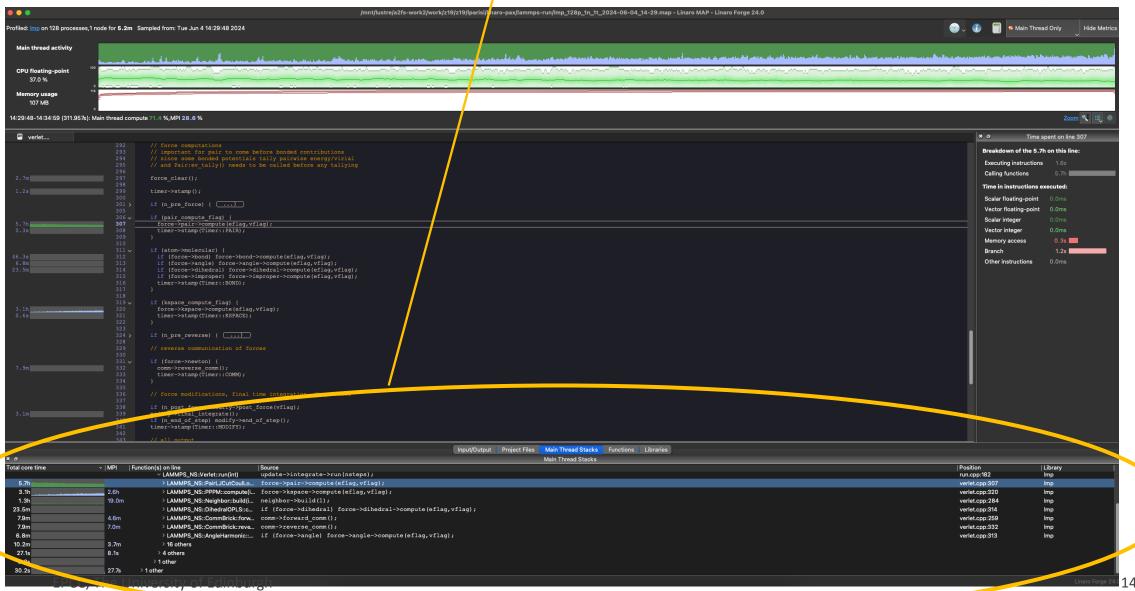
Sparklines next to the top timeconsuming lines





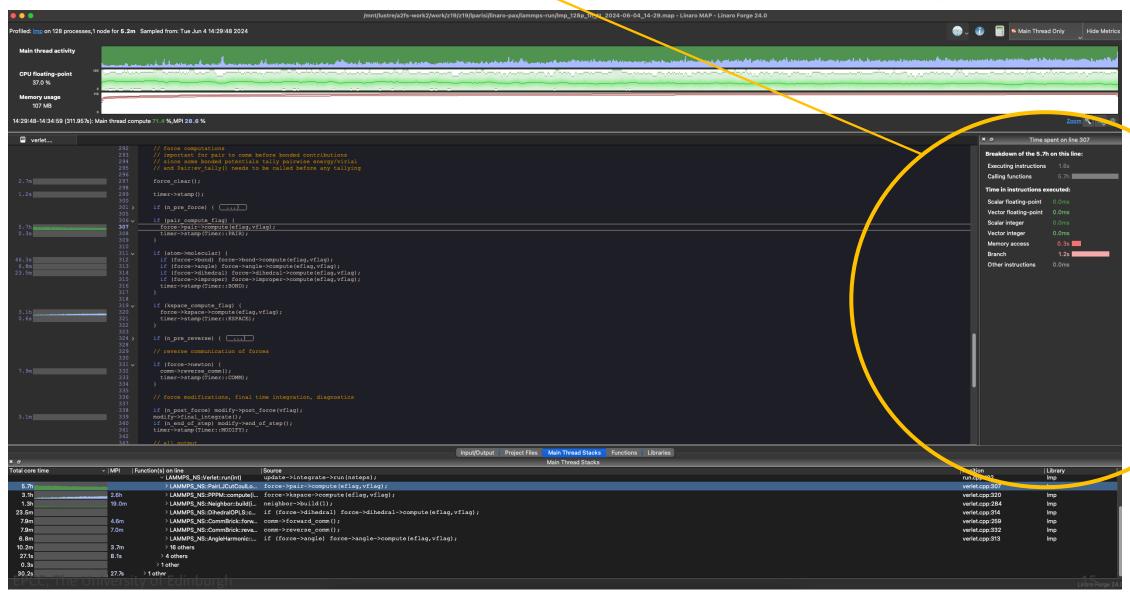
Stack of the main functions



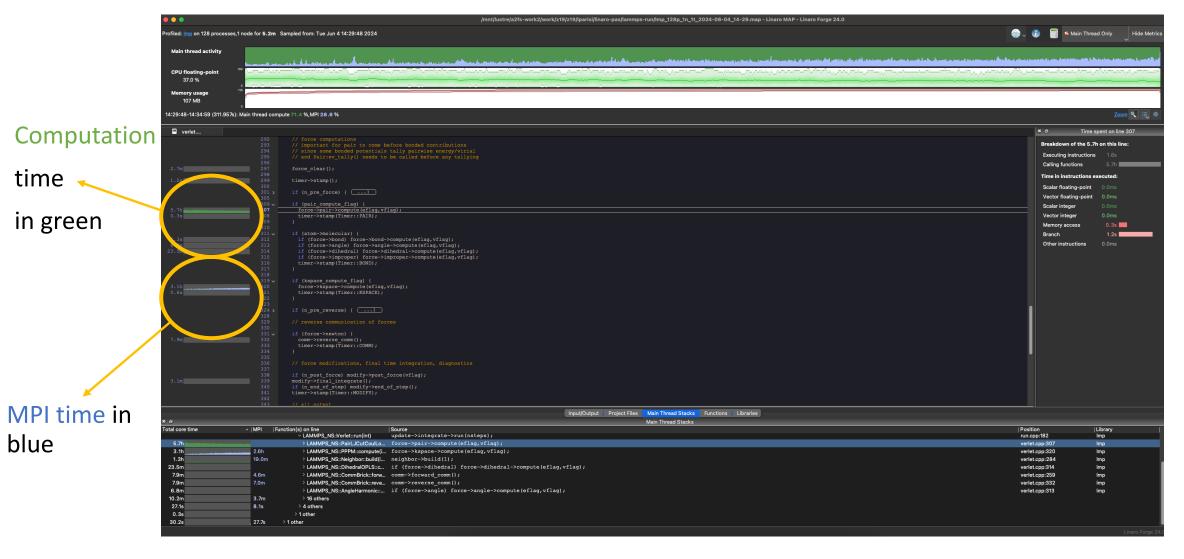


Most consuming instruction types On the selected line

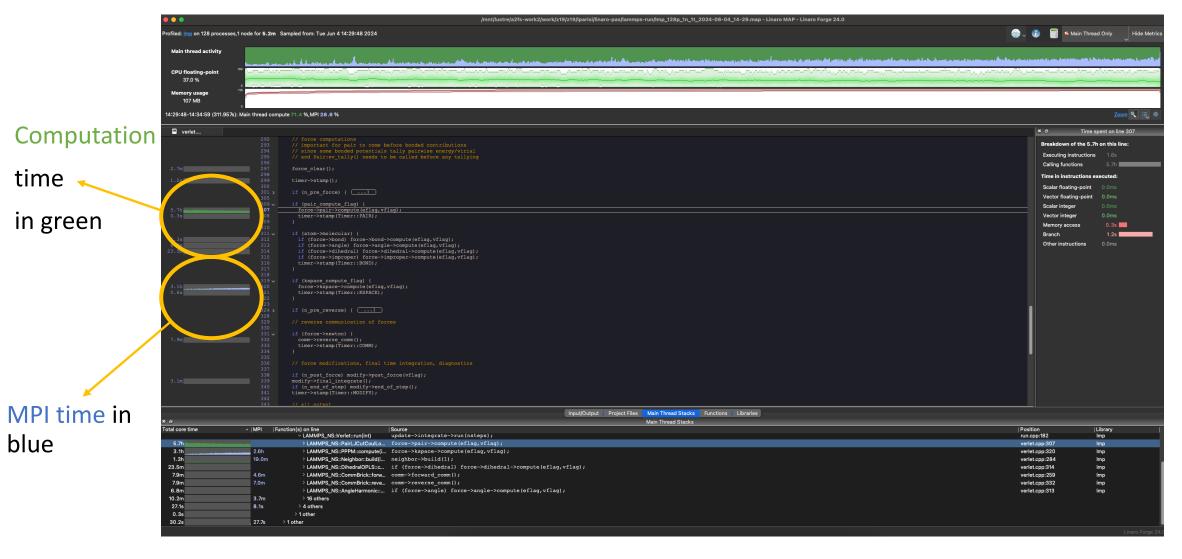










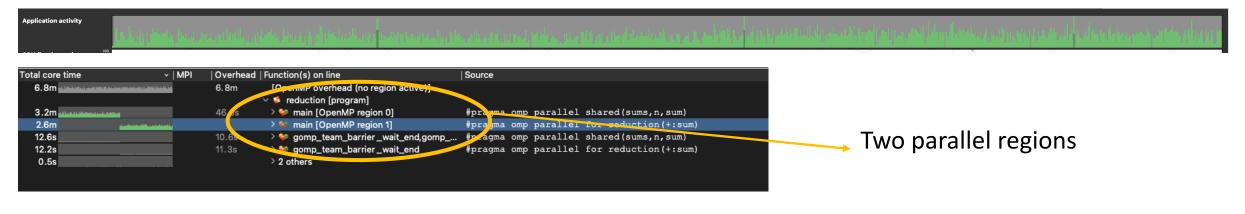


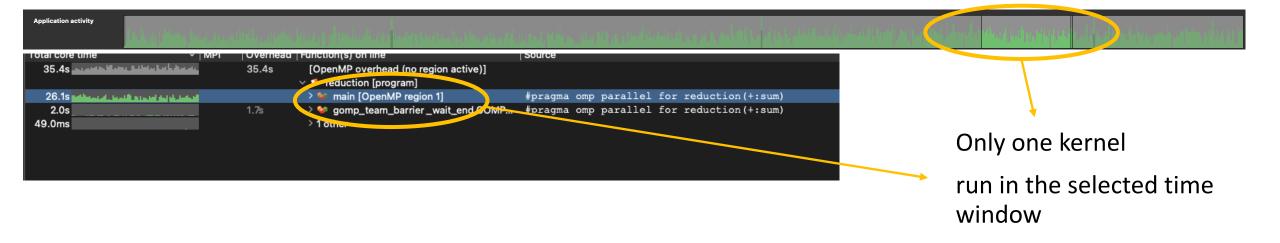


Hovering over the sparkline shows a quick breakdown of time consumption: MPI collectives, point-to-point, OpenMP overehad

Zooming in



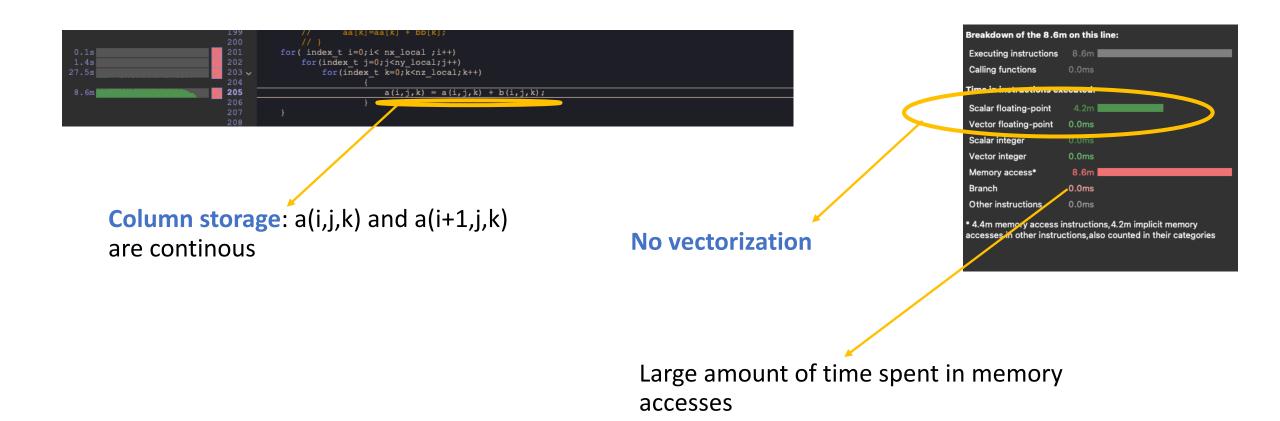






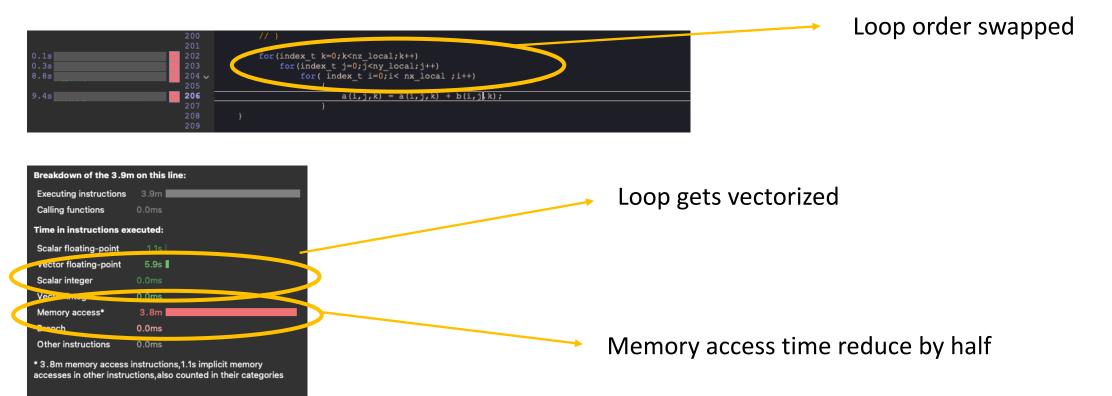
Memory access patterns





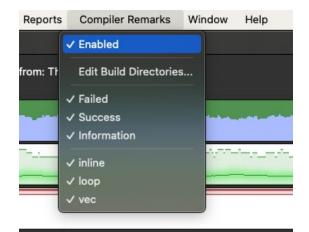
Memory access patterns





Compiler remarks

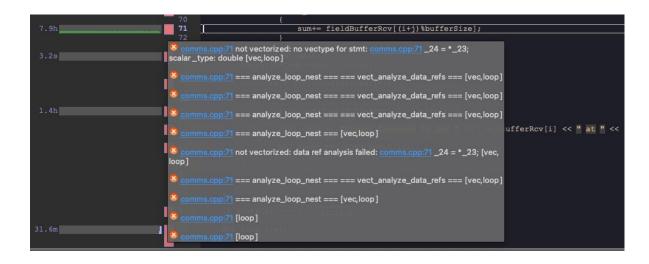




Compile with -fsave-optimization-record

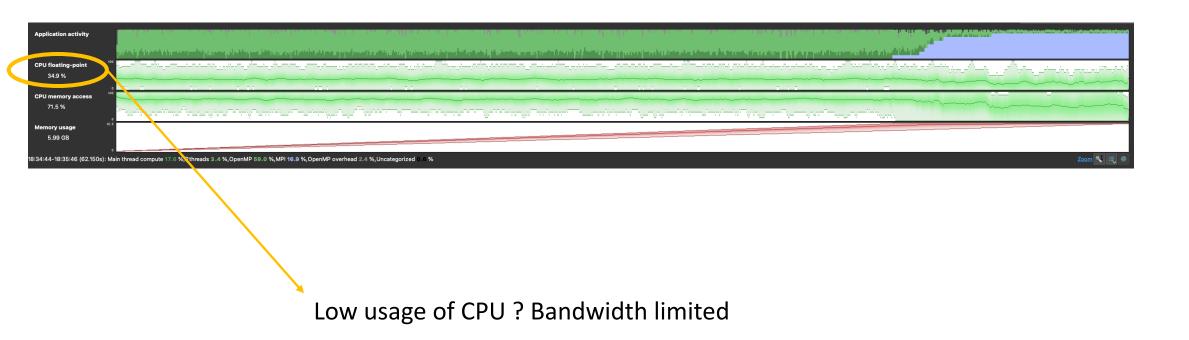
with GNU compiler

Check the Linaro Forge documentation for other compilers



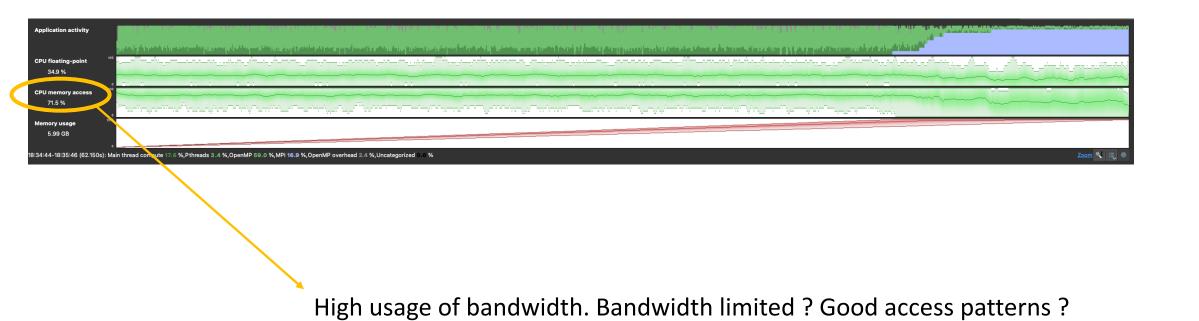
Memory usage





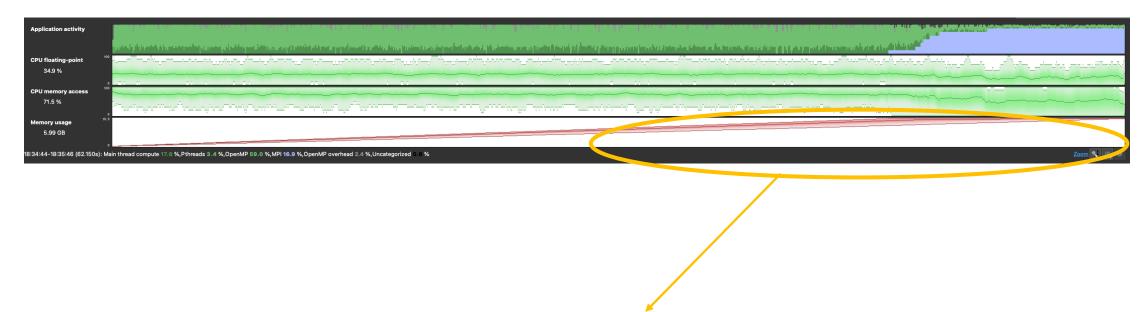
Memory usage





Memory leaks



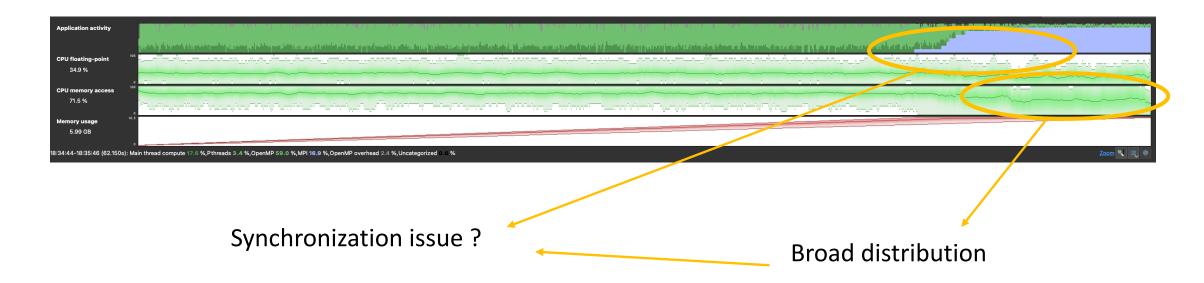


Memory keeps increasing. Memory leak?



Synchronization



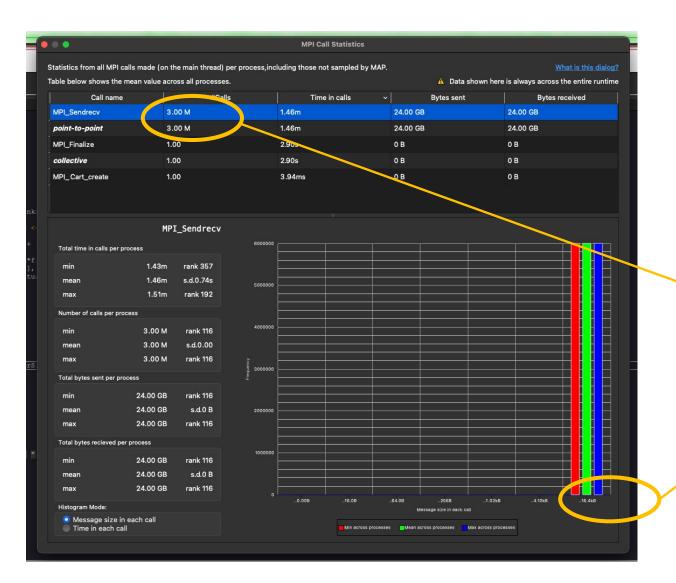


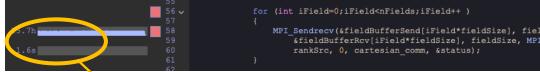


Time spent in an OpenMP barrier

MPI Communication







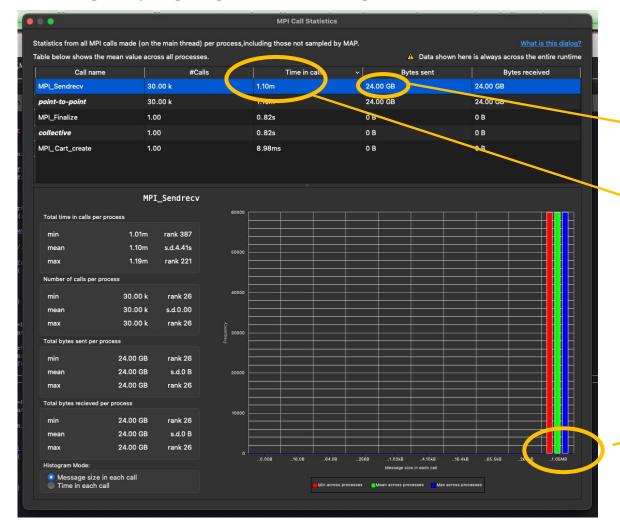
Large amount of CPU time spent in communication

Large number of MPI calls: 3M!!!

Small messages: about 16kb in size

MPI Communication

After grouping together messages...







Same total amount of data, but x100 less calls.

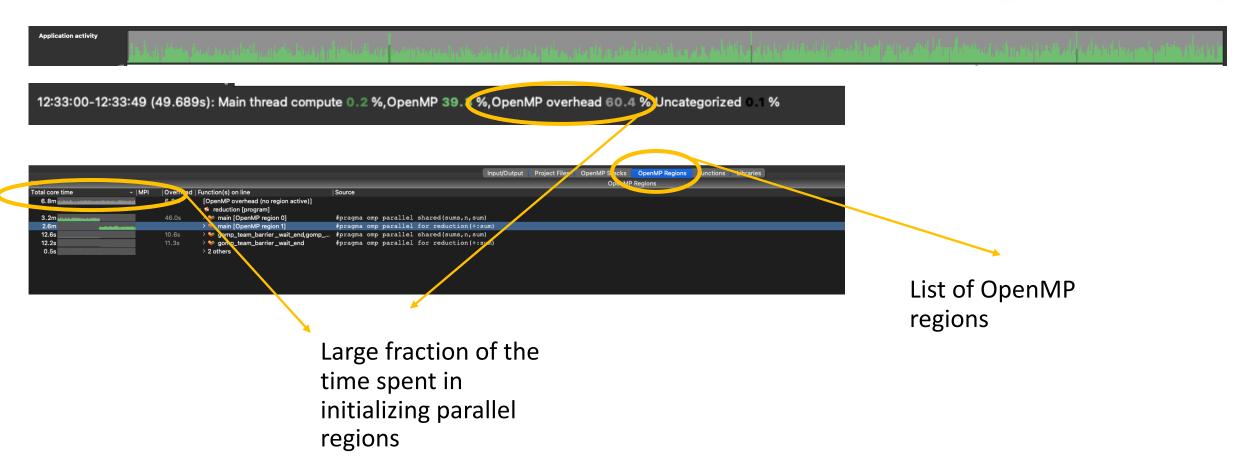
25% reduction in MPI wall time

Grouped many small messages in a large message



OpenMP

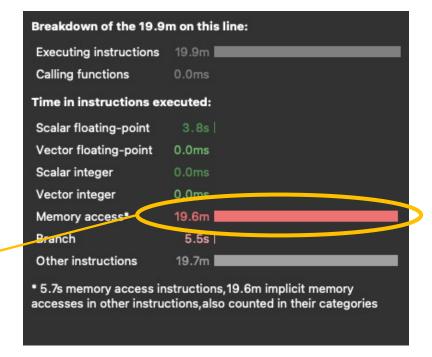




Atomics



Atomic operations show up as very long memory accesses





I/O bottlenecks



One rank spends a lot of time in IO

```
20 call initialise()

21 simulate()

23.59% of selected time (0.69s) in functions called from this line

23.6% of the selected time (0.38s) was in normal code on the main thread

0.93% of the selected time (15.00ms) was in point-to-point MPI calls on the nain thread

18.52% of the selected time (0.30s) was in I/O (reads & writes)

12.04% of the selected time (0.20s) was in I/O write and close calls

6.40 of the selected time (0.10s) was in I/O read calls
```

Large fraction of time spent writing to a file

Conclusions



- User friendly interface
- Quickly spot bottlenecks in your code
- Overview of communication patterns
- Overview of memory access and usage